
Handout #5: Project 1st Module

Due Feb 20th 2023, before 11:59 pm

Submit your report electronically via Gradescope.

General Instructions

- We will be making use of Jupyter notebooks to modularize and visualize the code. Please find the install instructions at the webpage:

<https://jupyter.org/install>

You may also find useful the following resources, which provide tutorials on how to use Jupyter notebooks and the python `numpy` library.

<https://www.codecademy.com/article/how-to-use-jupyter-notebooks>

<https://cs231n.github.io/python-numpy-tutorial/>

<https://numpy.org/doc/stable/user/quickstart.html>

- The project module directory structure is as follows:

```
Project_Module1
├── data
│   ├── regression
│   ├── binary_classification
│   └── fashion-mnist
├── codes
│   ├── Logistic.py
│   ├── MLogistic.py
│   ├── Regression.py
│   └── SVM.py
├── plots
├── utils
└── Notebook_Module1.ipynb
```

You will only need to make changes to the files in the `code` directory and the main `Notebook_Module1.ipynb` notebook. Please do not change the location of any files and maintain the same directory structure. Your generated plots would be stored in the `plots` folder. You do not have to write commands to load any data as that has already been done in the provided `Notebook_Module1.ipynb` file.

- **Important:** Any portions of the code that you must modify start with `YOUR CODE HERE` and end with `END YOUR CODE HERE`. Please do not change any code outside of these blocks.

- We will also make use of the following python libraries: `scipy`, `scikit-learn`, `matplotlib`, which can be installed using:

```
pip install -U scikit-learn scipy matplotlib
```

The code for loading and using these libraries has already been provided in the main file.

- You must upload a pdf of your report to Gradescope by **Monday, 20th February 11:59 p.m.** You may export your Jupyter notebook to a pdf format.
- Include a small section towards the end of your report on the contributions of each group member in the project.

1 Linear Regression

You will work through linear and polynomial regressions. Our data consists of inputs $x_n \in \mathbb{R}$ and targets $y_n \in \mathbb{R}$ for $n \in \{1, \dots, N\}$ which are related through a target function $y_n = f(x_n)$. Your goal is to learn a linear predictor $h_{\mathbf{w}}(x)$ that best approximates $f(x)$.

$$\mathbf{X} = \begin{bmatrix} x_1 \\ \vdots \\ x_N \end{bmatrix} \in \mathbb{R}^N, \quad \mathbf{y} = \begin{bmatrix} y_1 \\ \vdots \\ y_N \end{bmatrix} \in \mathbb{R}^N \quad (1)$$

The main file is the `Notebook_Module1.ipynb` Jupyter notebook.

- (a) (**Visualization**): Visualize the training and test data. What do you observe? For example, can you make an educated guess on the effectiveness of linear regression in predicting the data?
- (b) (**Linear Regression**): Recall that linear regression attempts to minimize the objective function

$$J(\mathbf{w}) = \frac{1}{N} \sum_{n=1}^N (h_{\mathbf{w}}(\mathbf{x}_n) - y_n)^2 \quad (2)$$

where $\mathbf{x}_n = (1, x_n)$. In this part we consider linear regression model, where $h_{\mathbf{w}}(\mathbf{x}_n) = \mathbf{w}^T \mathbf{x}_n = w_0 + w_1 x_n$. Note that to take into account the intercept term w_0 , we can add an additional *feature* to each instance and set it to one. This is equivalent to adding an additional first column to \mathbf{X} and setting it to all ones.

Modify the method `get_poly_features()` in `Regression.py` file for the case $m == 1$ to create a matrix \mathbf{X} for linear regression model.

- (c) Before tackling the harder problem of training the regression model, complete `predict()` in `Regression.py` file for the case $m = 1$ to predict \mathbf{y} from \mathbf{X} and \mathbf{w} .
- (d) Complete the function `loss_and_grad()` to compute the loss function and the gradient of the loss function with respect to \mathbf{w} for a data set \mathbf{X} and targets \mathbf{y} at given weights \mathbf{w} . Test your results by running the code in the main file `Notebook_Module1.ipynb`. If you implement everything correctly, you should get the loss function around 4 and gradient approximately $[-3.2, -10.5]$.
- (e) One way to solve linear regression is through gradient descent (GD). Complete the function `train_LR()` to train the linear regression model for given learning rate $\eta = 1e - 3$, `batch_size=30`, and number of iterations `num_iters=10000`. Plot the history of the loss function. What is the final value of the loss function and the value of the weight \mathbf{w} ? Experiment with different learning rates, batch sizes (e.g. full, stochastic), and number of iterations. What do you find works best for producing the best final loss function value?
- (f) We can also get a closed form expression to the linear regression. Complete `closed_form()` function in `Regression.py` to get the optimal weights \mathbf{w} . In the main file `Notebook_Module1.ipynb` to compare the optimal weights \mathbf{w} and loss function obtained from the closed form with the one obtained from GD.

- (g) (**Polynomial regression**) Now let us consider the more complicated case of polynomial regression where our hypothesis is:

$$h_{\mathbf{w}}(\mathbf{x}_n) = w_0 + w_1x_n + w_2x_n^2 + \dots + w_mx_n^m \quad (3)$$

Repeat steps (b) – (f) for the general case $m \geq 2$. Note that you only need to modify the `get_poly_features()` in `Regression.py` file for the case $m \geq 2$. Then, you will use the new features \mathbf{X} in the next steps. Plot the figures and values of the loss function for the case $m = 3$.

- (h) (**Overfitting**) For $m = \{1, \dots, 10\}$, in the main file `Notebook_Module1.ipynb`, complete the code by using the closed-form solver to determine the best polynomial regression model on the training data. With this model, calculate the loss on both the training data and the test data. Generate a plot depicting how the optimal loss varies with model complexity (polynomial degree). You should generate a single plot with both training and test error. Which degree polynomial would you say best fits the data? Was there evidence of under/overfitting the data? Use your plot to justify your answer.
- (i) (**Regularization**) For this problem, we will use ℓ_1 -regularization with the previous objective, which lead to the optimization objective:

$$J(\mathbf{w}) = \frac{1}{N} \sum_{n=1}^N (h_{\mathbf{w}}(\mathbf{x}_n) - y_n)^2 + \lambda \|\mathbf{w}\|_1 \quad (4)$$

Modify the method `loss_and_grad()` in the `Regression.py` file to reflect the effect of ℓ_1 -regularization for gradient based approach.

- (j) In the main file `Notebook_Module1.ipynb`, complete the code to find the coefficients that minimize the error for a third-degree polynomial ($m = 3$) given regularization factor $\lambda = \{10^{-8}, 10^{-7}, \dots, 10^{-2}\}$. Now use these coefficients to calculate the loss function on both the training data and test data as a function of λ using a learning rate of $\eta = 5e - 4$ and batch size of 10. Generate a plot depicting how the loss varies with λ (for your x-axis, let $x = \{1, 2, \dots, 7\}$ so that λ is on a logistic scale, with regularization increasing as x increases). Which λ value appears to work best?

2 Binary Classification

In this exercise, you will work through a family of binary classifications. Our data consists of inputs $x_n \in \mathbb{R}^{1 \times d}$ and labels $y_n \in \{-1, 1\}$ for $n \in \{1, \dots, N\}$. We will work on a subset of the Fashion-MNIST dataset which focuses on classifying whether the image is for a *Dress* ($y = 1$) or a *Shirt* ($y = -1$). Your goal is to learn a classifier based on linear predictor $h_{\mathbf{w}}(x) = \mathbf{w}^T x$. Let

$$\mathbf{X} = \begin{bmatrix} x_1 \\ \vdots \\ x_N \end{bmatrix} \in \mathbb{R}^{N \times d}, \quad \mathbf{y} = \begin{bmatrix} y_1 \\ \vdots \\ y_N \end{bmatrix} \in \{1, -1\}^N \quad (5)$$

The main file is the `Notebook_Module1.ipynb` Jupyter notebook.

- (a) (**Visualization**): Visualize a sample of the training data. What is the dimensions of X_{train} , and X_{test} .
- (b) (**Perceptron**): Implement Perceptron Algorithm to classify your training data. Let the maximum number of iterations of the Algorithm $num_{iter} = N$ (number of training samples). At each iteration, compute the percentage of misclassified points in the training dataset, and save it into a `Loss_hist` array. Plot the history of the loss function (`Loss_hist`). What is the final value of the loss function and the squared ℓ_2 norm value of the weight $\|\mathbf{w}\|_2^2$? Looking at the loss function, can you comment on whether the Perceptron algorithm converges?
- (c) (**Perceptron test error**): Compute the percentage of misclassified points in the test data for the trained Perceptron.
- (d) (**Logistic Regression**): In this part, we will implement the logistic regression for binary classification. Recall that logistic regression attempts to minimize the objective function

$$J(\mathbf{w}) = \frac{1}{N} \sum_{n=1}^N \log \left(1 + e^{h_{\mathbf{w}}(\mathbf{x}_n)} \right) - \sum_{n=1}^N \mathbf{1}_{y_n=1} h_{\mathbf{w}}(\mathbf{x}_n) \quad (6)$$

where $\mathbf{x}_n = (1, x_n)$, and $\mathbf{1}_A = 1$ if A is true and 0 otherwise. Moreover, $h_{\mathbf{w}}(\mathbf{x}_n) = \mathbf{w}^T \mathbf{x}_n$. First, we will add an additional *feature* to each instance and set it to one. This is equivalent to adding an additional first column to \mathbf{X} and setting it to all ones.

Modify the `get_features()` in `Logistic.py` file to create a matrix \mathbf{X} for logistic regression model [You might use the function `get_poly_features()` from the previous question].

- (e) Complete `predict()` in `Logistic.py` file to predict \mathbf{y} from \mathbf{X} and \mathbf{w} .
- (f) Complete the function `loss_and_grad()` to compute the loss function and the gradient of the loss function with respect to \mathbf{w} for a data set \mathbf{X} and labels \mathbf{y} at given weights \mathbf{w} . Test your results by running the code in the main file `Notebook_Module1.ipynb`. If you implement everything correctly, you should get the loss function within 0.7 and squared ℓ_2 norm of the gradient around 1.8×10^5 .
- (g) Complete the function `train_LR()` to train the logistic regression model for given learning rate $\eta = 10^{-6}$, `batch_size` = 100, and number of iterations $num_{iters} = 5000$. Plot the history of the loss function (`Loss_hist`). What is the final value of the loss function and the squared ℓ_2 norm value of the weight $\|\mathbf{w}\|_2^2$?
- (h) (**Logistic Regression test error**): Compute the percentage of misclassified points in the test data for the trained Logistic Regression.
- (i) (**Logistic Regression and Batch Size**): Train the Logistic regression model with different batch size $b \in \{1, 50, 100, 200, 300\}$, learning rate $\eta = 10^{-5}$, and number of iterations $num_{iter} = 6000/b$. Train each model 10 times and average the test error for each value of batch size. Plot the test error as a function of the batch size. Which batch size gives the minimum test error?
- (j) (**SVM**): In this part, we will implement SVM for binary classification. Recall that SVM attempts to minimize the objective function

$$J(\mathbf{w}) = \frac{1}{N} \sum_{n=1}^N \max(0, 1 - y_n h_{\mathbf{w}}(\mathbf{x}_n)) \quad (7)$$

where $\mathbf{x}_n = (1, x_n)$, and $\mathbf{1}_A = 1$ if A is true and 0 otherwise. Moreover, $h_{\mathbf{w}}(\mathbf{x}_n) = \mathbf{w}^T \mathbf{x}_n$. First, we will add an additional *feature* to each instance and set it to one. This is equivalent to adding an additional first column to \mathbf{X} and setting it to all ones.

Modify the `get_features()` in `SVM.py` file to create a matrix \mathbf{X} for SVM model.

- (k) Complete `predict()` in `SVM.py` file to predict \mathbf{y} from \mathbf{X} and \mathbf{w} .
- (l) Complete the function `loss_and_grad()` in `SVM.py` to compute the loss function and the gradient of the loss function with respect to \mathbf{w} for a data set \mathbf{X} and labels \mathbf{y} at given weights \mathbf{w} . Test your results by running the code in the main file `Notebook.Module1.ipynb`. If you implement everything correctly, you should get the loss function around 1 and squared ℓ_2 norm of gradient approximately 7.5×10^5 .
- (m) Complete the function `train_svm()` to train the svm model for given learning rate $\eta = 10^{-6}$, `batch_size` = 50, and number of iterations `num_iters` = 5000. Plot the history of the loss function (`Loss_hist`). What is the final value of the loss function and the value of squared ℓ_2 norm of the weight $\|\mathbf{w}\|_2^2$?
- (n) (**SVM test error**): Compute the percentage of misclassified points in the test data for the trained SVM.
- (o) (**SVM and Batch Size**): Train the SVM model with different batch size $b \in \{1, 50, 100, 200, 300\}$, learning rate $\eta = 10^{-5}$, and number of iterations `num_iter` = $6000/b$. Train each model 10 times and average the test error for each value of batch size. Plot the test error as a function of the batch size. Which batch size gives the minimum test error?
- (p) (**Kernelized SVM**): In this part we will use `sklearn.svm.SVC` library to train a RBF kernelized SVM. RBF kernelized SVM uses a parameter C that represents an inverse ℓ_2 -regularization strength (higher C corresponds to a lower regularization). Given the list of possible C values $\{0.01, 0.1, 0.25, 0.75, 1\}$ find the best value by plotting test error vs. C . How is the test error compared to non-kernelized SVM?

3 Multi-Class Classification

In this exercise, you will consider a multi-class image classification task for the Fashion-MNIST dataset [2]. Our data consists of inputs $x_n \in \mathbb{R}^{1 \times d}$ and labels $y_n \in \{0, 1, \dots, 9\}$ for $n \in [N]$. Thus:

$$\mathbf{X} = \begin{bmatrix} x_1 \\ \vdots \\ x_N \end{bmatrix} \in \mathbb{R}^{N \times d}, \quad \mathbf{y} = \begin{bmatrix} y_1 \\ \vdots \\ y_N \end{bmatrix} \in \{0, \dots, 9\}^N \quad (8)$$

The main file is `Notebook.Module1.ipynb` Jupyter notebook.

- (a) (**Visualization**): Visualize a sample of the training data. What is the dimensions of X_{train} , and X_{test} ?
- (b) (**Multi-class Logistic Regression**): In this part, we will implement the Multi-class logistic regression. Recall that logistic regression attempts to minimize the objective function

$$J(\mathbf{W}) = \frac{1}{N} \left[\sum_{i=1}^N \log \left(\sum_{j=1}^k e^{w_j^T \mathbf{x}_i} \right) - \sum_{i=1}^N w_{y_i}^T \mathbf{x}_i \right] \quad (9)$$

where $\mathbf{x}_n = (1, x_n)$, and $\mathbf{W} = [w_1, \dots, w_k] \in \mathbb{R}^{k \times d}$ weight matrix, where $w_i \in \mathbb{R}^d$. First, we will add an additional *feature* to each instance and set it to one. This is equivalent to adding an additional first column to \mathbf{X} and setting it to all ones.

Modify the `get_features()` in `MLogistic.py` file to create a matrix \mathbf{X} for logistic regression model (You may use the function `get_poly_features()` from the previous questions).

- (c) Complete `predict()` in `MLogistic.py` file to predict \mathbf{y} from \mathbf{X} and \mathbf{W} . Recall that for a new feature \mathbf{x} , we assign \mathbf{x} to a class j^* such that

$$j^* = \arg \max_{j \in \{1, \dots, k\}} \frac{e^{w_j^T \mathbf{x}}}{\sum_{i=1}^k e^{w_i^T \mathbf{x}}} \quad (10)$$

- (d) Complete the function `loss_and_grad()` to compute the loss function and the gradient of the loss function with respect to \mathbf{W} for a dataset \mathbf{X} and labels \mathbf{y} at given weights \mathbf{W} . Test your results by running the code in the main file `Notebook_Module1.ipynb`. If you implement everything correctly, you should get the loss function around 2.3 and squared Frobenius norm of gradient matrix approximately 420. Note that \mathbf{W} is a $k \times d$ matrix. Hence, the dimension of the gradient of the function J with respect to \mathbf{W} is a matrix with the same size of $k \times d$.
- (e) Complete the function `train_LR()` to train the logistic regression model for given learning rate $\eta = 1e-7$, `batch_size` = 200, and number of iterations `num_iters` = 1500. Plot the history of the loss function (`Loss_hist`). What is the final value of the loss function and the square Frobenius norm of the weight matrix $\|\mathbf{W}\|_F^2$?
- (f) (**Logistic Regression test error**): Compute the percentage of the misclassified points in the test data for the trained Logistic Regression.
- (g) (ℓ_1 **Regularization**): In this part, we will use ℓ_1 -regularization so that our regularized objective function is

$$J(\mathbf{W}) = \frac{1}{N} \left[\sum_{i=1}^N \log \left(\sum_{j=1}^k e^{w_j^T \mathbf{x}_i} \right) - \sum_{i=1}^N w_{y_i}^T \mathbf{x}_i \right] + \lambda \sum_{j=1}^k \sum_{l=1}^d |w_{jl}| \quad (11)$$

Modify the function `loss_and_grad()` by adding the effect of regularization.

- (h) Consider the allowed set of values of λ as $\Lambda := \{0, 10^{-6}, 10^{-3}, 10^{-2}, 10^{-1}, 1\}$.
- In the main file `Notebook_Module1.ipynb`, complete the code to train the logistic regression model with regularization parameter $\lambda \in \Lambda$. Now use the obtained final model to find the test error and the final training error as a function of λ . You must train each model using learning rate $1e^{-7}$, `batch_size`=200 and number of iterations 3000. Generate a plot depicting how the test error varies with λ (x-axis represents λ). Which λ value appears to work best?
 - Now use k -fold cross validation (with $k = 5$) to find a suitable value for the regularization coefficient. Specifically, divide the training dataset (at random) into $k = 5$ parts and at each round, pick one of these parts as a validation dataset and the other remaining parts as the new training dataset. We can now calculate the validation error for a model trained on the new train dataset. The final validation error for a given λ is the average validation error across the $k = 5$ parts. Compute the average cross validation error for

Algorithm 1 Multi-Class Adaboost

:

- 1: **Inputs:** N training samples \mathbf{X} , \mathbf{y} , Total number of weak classifiers T .
- 2: **Initialization:** Sample weights $\mathbf{D} = [d_1, \dots, d_N]$ s.t. $d_i = 1/N$ for all $i \in \{1, \dots, N\}$.
- 3: **for** $m = 1$ to T **do**
- 4: Train a decision tree classifier h_m with maximum depth 4 using training samples (\mathbf{X}, \mathbf{y}) and $sample_weight = \mathbf{D}$
- 5: Compute $err[m] \leftarrow \sum_{i=1}^N d_i \mathbb{I}(y_i \neq h_m(\mathbf{x}_i))$
- 6: Compute $\alpha_m \leftarrow \log\left(\frac{1-err[m]}{err[m]}\right) + \log(k-1)$
- 7: Update sample weights: $d_i \leftarrow d_i \exp(\alpha_m \mathbb{I}(y_i \neq h_m(\mathbf{x}_i))) \quad \forall i \in \{1, \dots, N\}$
- 8: Normalize the sample weights $d_i = \frac{d_i}{\sum_{j=1}^N d_j} \quad \forall i \in \{1, \dots, N\}$
- 9: **end for**
- 10: **Outputs:** For a new feature \mathbf{x} , assign a class:

$$J^*(\mathbf{x}) = \arg \max_{j \in \{0, \dots, k-1\}} \sum_{m=1}^T \alpha_m \mathbb{I}(h_m(\mathbf{x}) = j)$$

each $\lambda \in \Lambda$. What λ value seems to give the smallest average validation error? Find the test error for the model trained (on the entire train dataset) using this value of λ . As before, train each model using learning rate $1e^{-7}$, `batch_size=200` and number of iterations 3000.

- (i) (**Decision Trees and Adaboost**): In this part, we will implement Adaboost algorithm for Multi-class classification [1]. First, install `scikit-learn`, using the following command:

```
pip install -U scikit-learn scipy matplotlib
```

Implement the Adaboost algorithm 1. To train a decision tree classifier, you can use the following command:

```
tree = DecisionTreeClassifier(max_depth = 4).fit(X_train, y_train,  
sample_weight = D)
```

where maximum depth $max_depth = 4$ and sample weights $sample_weight = D$. To predict the labels from the features, you can use:

```
tree.predict(X_test)
```

Plot a figure depicting the test error and the training error in the same plot as a function of the number of classifiers (T). What is the test error for $T = 1$ and the test error for $T = 200$?

References

- [1] Hastie, T., Rosset, S., Zhu, J. and Zou, H., 2009. Multi-class adaboost. *Statistics and its Interface*, 2(3), pp.349-360.
- [2] <https://github.com/zalandoresearch/fashion-mnist>